

The SpaceX logo is displayed in a stylized, white, sans-serif font. A bright orange arc, representing a rocket's trajectory, curves from the bottom left towards the top right, passing behind the letters of the logo.

# SPACEX



## IBM Data Science Professional Certificate Capstone Project

Ingrid Lindstrom  
February 2022

# Outline

---

- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

# Executive Summary

---

- **Methodologies Utilized**

- Data Collection through API
- Data Collection with Web Scraping
- Data Wrangling
- Exploratory Data Analysis with SQL
- EDA with Data Visualization
- Interactive Data Analytics with Folium
- Machine Learning Prediction

- **Results Summary**

- Exploratory Data Analysis results
- Interactive Analytics demonstration with screenshots
- Predictive Analytics results

# Introduction

---

- **Project Background and Context**

The purpose of this project is to predict if the Falcon 9 first stage will land successfully. SpaceX advertises Falcon 9 rocket launches on its website, with a cost of \$62 million dollars. Other providers cost upward of \$165 million dollars each. The majority of these savings is due to SpaceX's ability to reuse the first stage of the rocket. If I can determine if the first stage of a Falcon 9 rocket will land, I can determine the cost of the launch. This financial information can be used if a rival competitor wants to bid against SpaceX for rocket launches and other aerospace engineering endeavors.

- **Pertinent Questions that can be answered with Data Science techniques**

- What **factors** determine if the rocket will land successfully?
- Which parameters determine the **success rate** of a successful landing?
- What **operating conditions** need to be in place to ensure a successful landing?



# Section 1

---

## Methodology



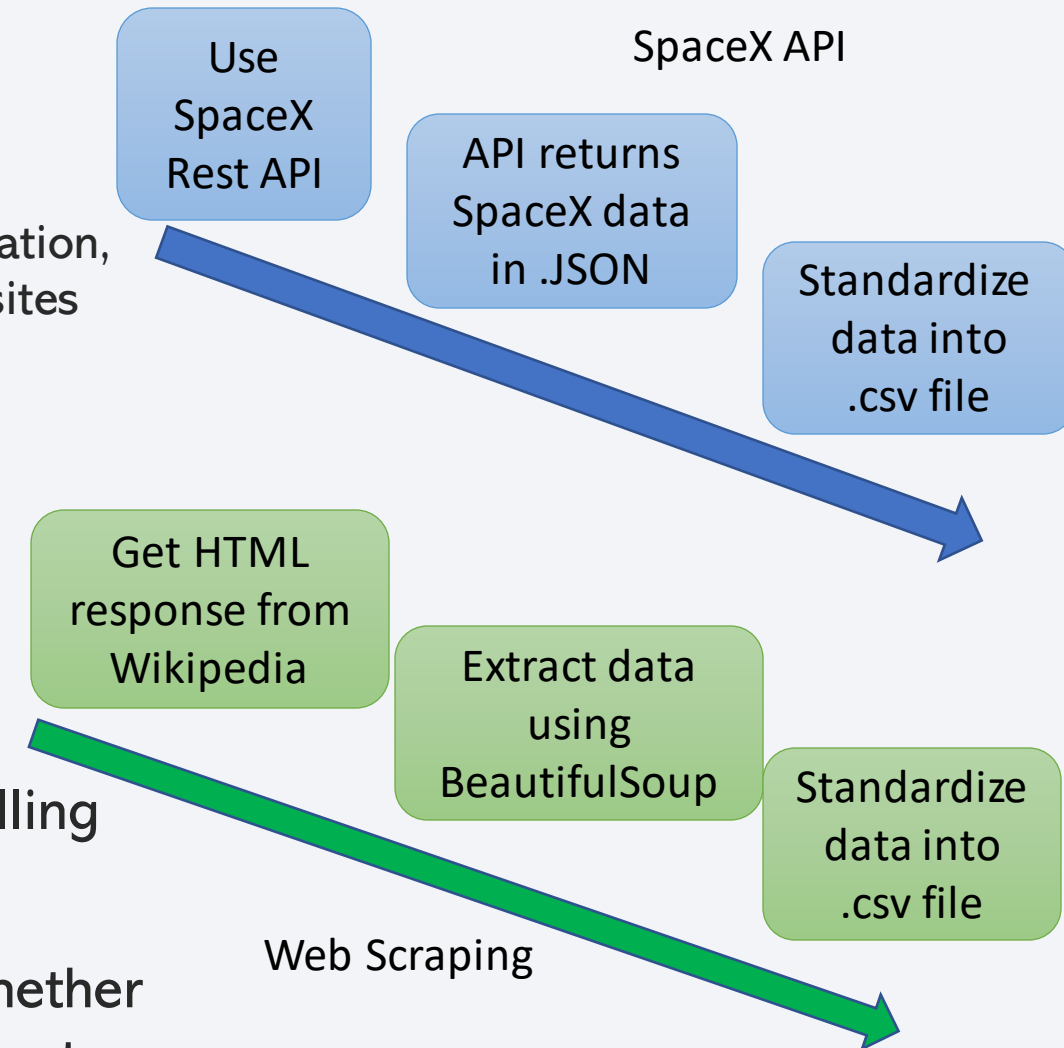
# Methodology Overview

---

- Data collection methodology
  - SpaceX Rest API
  - Web Scraping from Wikipedia – using Python library *BeautifulSoup*
- Perform data wrangling
  - One Hot Encoding was applied to categorical features in the dataframe
  - OHE prepares fields for Machine Learning, eliminating redundant or irrelevant columns
- Perform exploratory data analysis (EDA) using visualization and SQL
- Perform interactive visual analytics using Folium and Plotly Dash – **dynamic results in real time!**
- Perform predictive analysis using classification models
  - Building, tuning, and evaluating the highest accuracy of various classification models

# Data Collection

- The primary data set used in this project was gathered from the SpaceX Rest API.
  - This API gives us the relevant data about launch information, including the model number of the rocket, geographic sites utilized, the payload delivered, launch and landing specifications, and landing outcome.
  - The SpaceX Rest endpoints, or URL, is obtained from [api.spacexdata.com/v4/](https://api.spacexdata.com/v4/)
- An additional data source for obtaining Falcon 9 launch data is web scraping a Wikipedia page using *BeautifulSoup*, a Python library used for pulling data out of XML and HTML files.
- The ultimate goal is to use this data to predict whether SpaceX will attempt to land a Falcon 9 rocket or not.





# Data Collection – SpaceX API



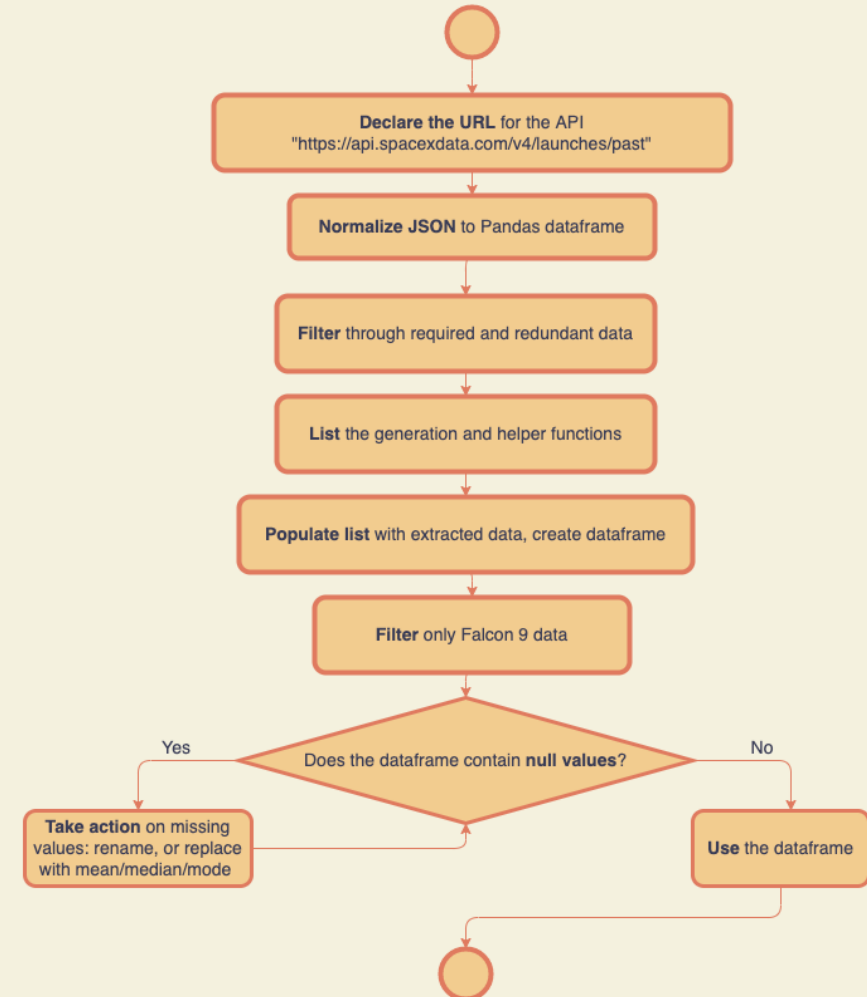
**Core Procedure:** Declare URL, Normalize JSON, Filter, List, Populate, Specific Filter, Take Action on Null Values



GitHub URL to [API Jupyter Notebook](#)

## SpaceX Data Collection with API Procedure

Ingrid Lindstrom







# Data Collection – Web Scraping



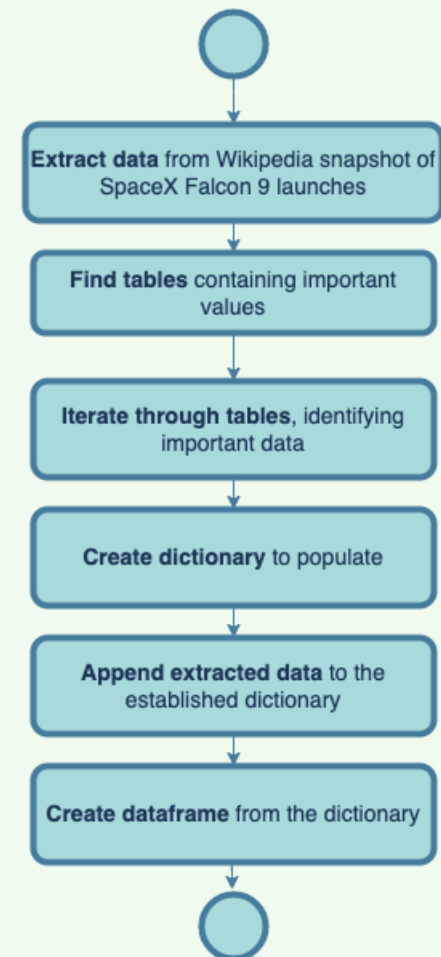
**Core Procedure:** Extract Data, Find Tables, Iterate, Create Dictionary, Append Extracted Data, Create Dataframe



GitHub URL link to [Web Scraping Jupyter Notebook](#)


## SpaceX Web Scraping Procedure

Ingrid Lindstrom





# Data Wrangling



Each SpaceX launch is aimed at a dedicated orbit. This figure depicts some common orbit types.



In the data set, there are several possible outcomes for booster landings; these are labeled appropriately for sorting and analysis.



I converted these outcomes to binary training labels: 1 for successful or 0 for unsuccessful



True Ocean: successful landing to a specific ocean region; False Ocean: unsuccessful landing



True RTLS: successful landing to on a ground pad; False RTLS: unsuccessful landing



True ASDS: successful landing on a drone ship; False ASDS: unsuccessful landing



GitHub URL link to [Data Wrangling Notebook](#)

# Data Wrangling, continued

---

This is the procedure utilized for converting outcomes into training labels to determine 1 for a successful booster landing, or 0 for an unsuccessful booster landing.

Perform Exploratory Data Analysis (EDA) on dataset

Calculate the number of **launches** at each site

Calculate the **number** and **occurrence** of each orbit

Calculate the number and occurrence of **mission outcome per orbit type**

Work out the success rate for every landing in the data set

Create a "Landing Outcome" label from the Outcome column

Export dataset as .csv file


# EDA with Data Visualization

[GitHub Link to EDA Notebook](#)

## Scatterplot Graphs

Used to show the relationship between two variables; best for large datasets.


- Flight Number vs. Payload Mass
- Flight Number vs. Launch Site
- Payload vs. Launch Site
- Orbit vs. Flight Number
- Payload vs. Orbit Type
- Orbit vs. Payload Mass



## Bar Graphs

Convenient to compare categories on one axis, and discrete values on the other. Used to see clear differences between groups.

- Mean Value vs. Orbit




## Line Graphs

These depict variables and trends over time; extrapolations are used to form predictions.

- Success Rate vs. Year





## DISPLAY QUERIES

- ✓ Display all unique names of distinct launch sites for each rocket launch mission
- ✓ Display 5 records where the launch sites begin with the string "KSC"
- ✓ Display the **total** payload mass (kg) carried by boosters launched by NASA (CRS)
- ✓ Display the average payload mass (kg) carried by Falcon 9 version 1.1



## LISTING QUERIES

- ✓ List the date of the successful landing outcome [drone ship]
- ✓ List the names of the boosters of successful landing outcomes [ground pad] with subquery of calculating payload mass greater than 4000 but less than 6000 kg
- ✓ List the **total number** of successful and failure outcomes
- ✓ List the names of [Booster\_Versions] which carried the maximum payload mass
- ✓ List the records of month names, successful landing outcomes for [ground pad], booster versions, launch sites in 2017




## RANKING QUERY

- ✓ Rank the count of successful landing outcomes between June 2010 and March 2017 in descending order

# Build an Interactive Map with Folium

- Folium allows us to illustrate data with movement, and easily identifies successful or failed launches
  - Take Latitude & Longitude coordinates of each launch site
  - Add a circle marker around each site, labeled clear with the site name
- Assign dataframe launch\_outcomes(failures, successes) to class 0 and 1 with **red** and **green** markers, respectively, using MarkerCluster()
- Calculate the distance between Launch Sites and various landmarks (railways, highways, coastline) to identify trends. Lines are drawn on the map to indicate distance



## Launch Site Trends

Close in proximity to railways? **NO**

Close in proximity to highways? **NO**

Close in proximity to coastlines? **YES**

Certain distance away from cities? **YES**

GitHub URL Link to [Interactive Visual Analytics with Folium Jupyter Notebook](#)

# Build a Dashboard with Plotly Dash

---

A live website dashboard built with Dash shows real-time data updates with graphs and charts.

CHARTS	GRAPHS
<ul style="list-style-type: none"><li>• A <b>pie chart</b> shows the total launches at a specific site or ALL sites</li><li>• The size of the circle (pie) can be scaled proportionally to the total quantity</li><li>• Plotly Dash displays relative proportions of multiple classes of data</li></ul>	<ul style="list-style-type: none"><li>• A <b>scatterplot</b> shows the relationship between Launch Outcome and Payload Mass (kg) for various Booster Versions</li><li>• The relationship between two variables are depicted</li><li>• Most reliable method to display a non-linear pattern</li><li>• Minimum and maximum values are shown in a range</li></ul>

# Predictive Analysis (Classification)

GitHub Link to [Machine Learning Prediction Jupyter Notebook](#)

## 1. BUILD

- Load the dataset into NumPy & Pandas
- Transform the data
- Split the data into training and testing sets
- Consider which algorithms should be used
- Set the parameters & algorithms to GridSearchCV
- Fit the datasets into GridSearchCV objects and train the models

## 4. DECIDE

- Look at the hard numbers
- The model with the best accuracy score = **best performance**

## 2. EVALUATE

- Check the accuracy of each model
- Tune the hyperparameters for each type of algorithm
- Plot the Confusion Matrix, make sure colors are consistent

## 3. IMPROVE

- Proceed to feature engineering
- Algorithm tuning
- Re-run the accurate models, noting any discrepancies





# Results

---

- Exploratory Data Analysis results
  - SQL comparisons, contrasts, lists, rankings, and selections
- Interactive Analytics demo in screenshots
  - Folium, Dash charts and scatterplots
- Predictive Analysis results
  - Examining the most effective machine learning models to  
*turn data into decisions!*




# Section 2

---

Insights Drawn from  
Exploratory Data Analysis (EDA)

# Flight Number vs. Launch Site


---



If the number of flights is greater than 30, the higher the success rate of the launch at that site.

# Payload vs. Launch Site


---



At launch site CCAFS SLC 40, the greater the payload mass, the higher the success rate for the rocket launch.


# Success Rate vs. Orbit Type

---




Four orbits: ES-L1, GEO, HEO, and SSO, had the highest mean success rates. These are the most stable and dependable orbits for the SpaceX rockets.

# Flight Number vs. Orbit Type



- In LEO orbits, successful outcomes are correlated with the number of flights
- In GTO orbits, there is no relationship between number of flights and successful outcomes
- Similarly, SSO orbits are consistently successful, but does not appear to be dependent on number of flights


# Payload vs. Orbit Type



In launches with heavier payloads, successful landings are higher for PO, LEO and ISS orbits.

# Launch Success Yearly Trend

---



Except for a dip in 2018, the launch success rate has continuously increased each year from 2013 through 2020.



# All Launch Site Names

---

Display the names of the unique launch sites in the space mission

```
In [7]: df["Launch_Site"].unique()
```

```
Out[7]: array(['CCAFS LC-40', 'VAFB SLC-4E', 'KSC LC-39A', 'CCAFS SLC-40'],  
          dtype=object)
```

- SQLite Query Word: UNIQUE
- This command shows only unique values in the **Launch\_Site** column from **tblSpaceX**

# Launch Site Names Begin with "KSC"

Display 5 records where launch sites begin with the string 'KSC'

```
In [8]: df["Launch_Site"].str.startswith("KSC")
b = df["Launch_Site"].str.startswith("KSC")
df[b].head()
```

```
Out[8]:
```

	Date	Time (UTC)	Booster_Version	Launch_Site	Payload	PAYLOAD_MASS_KG_	Orbit	Customer	Mission_Outcome	Landing_Outcome
29	19-02-2017	14:39:00	F9 FT B1031.1	KSC LC-39A	SpaceX CRS-10	2490	LEO (ISS)	NASA (CRS)	Success	Success (ground pad)
30	16-03-2017	06:00:00	F9 FT B1030	KSC LC-39A	EchoStar 23	5600	GTO	EchoStar	Success	No attempt
31	30-03-2017	22:27:00	F9 FT B1021.2	KSC LC-39A	SES-10	5300	GTO	SES	Success	Success (drone ship)
32	01-05-2017	11:15:00	F9 FT B1032.1	KSC LC-39A	NROL-76	5300	LEO	NRO	Success	Success (ground pad)
33	15-05-2017	23:21:00	F9 FT B1034	KSC LC-39A	Inmarsat-5 F4	6070	GTO	Inmarsat	Success	No attempt

- SQLite Query Words: STARTSWITH
- Using **HEAD** will show only five records from the dataframe; **STARTSWITH** will yield only **Launch\_Site** records beginning with the string "KSC"

# Total Payload Mass

---

Display the total payload mass carried by boosters launched by NASA (CRS)

```
In [9]: df["Customer"] == "NASA (CRS)"  
n = df["Customer"] == "NASA (CRS)"  
df[n]["PAYLOAD_MASS__KG_"].sum()
```

```
Out[9]: 45596
```

- SQLite Query Word: SUM
- The total payload mass carried by NASA CRS boosters was 45,496kg.

# Average Payload Mass by F9 v1.1

---

Display average payload mass carried by booster version F9 v1.1

```
In [12]: df["Booster_Version"] == "F9 v1.1"  
F9 = df["Booster_Version"] == "F9 v1.1"  
df[F9]["PAYLOAD_MASS__KG_"].mean()
```

```
Out[12]: 2928.4
```

- SQLite Query Word: MEAN
- The average payload mass carried by booster version F9 v1.1 is 2,928.4kg

# First Successful Ground Landing Date

List the date where the succesful landing outcome on the drone ship was acheived.

*Hint: Use min function*

```
In [25]: df["Landing_Outcome"] == "Success (drone ship)"
Achieve = df["Landing_Outcome"] == "Success (drone ship)"
df[Achieve]
df[Achieve]["Date"]
```

```
Out[25]: 22    08-04-2016
23    06-05-2016
24    27-05-2016
27    14-08-2016
28    14-01-2017
31    30-03-2017
35    23-06-2017
36    25-06-2017
39    24-08-2017
41    09-10-2017
42    11-10-2017
43    30-10-2017
52    18-04-2018
53    11-05-2018
Name: Date, dtype: object
```

- SQLite Query Words: DATE, MIN
- All the dates in the list correspond to a successful [drone ship] landing outcome. Using MIN would reveal that the earliest date is April 8, 2016.

# Successful Drone Ship Landing with Payload between 4000 and 6000

```
List the names of the boosters which have success in ground pad and have payload mass greater than 4000 but less than 6000

In [30]: df["Landing_Outcome"].unique()

Out[30]: array(['Failure (parachute)', 'No attempt', 'Uncontrolled (ocean)',
              'Controlled (ocean)', 'Failure (drone ship)',
              'Precluded (drone ship)', 'Success (ground pad)',
              'Success (drone ship)', 'Success', 'Failure', 'No attempt'],
              dtype=object)

In [39]: (df["PAYLOAD_MASS_KG_"] < 4000) & (df["PAYLOAD_MASS_KG_"] > 6000)

Out[39]: 0      False
          1      False
          2      False
          3      False
          4      False
          ...
          96     False
          97     False
          98     False
          99     False
         100     False
          Name: PAYLOAD_MASS_KG_, Length: 101, dtype: bool
```

```
In [52]: Boostername = (df["Landing_Outcome"] == "Success (ground pad)") & (df["PAYLOAD_MASS_KG_"] > 4000) & (df["PAYLOAD_MASS_KG_"] < 6000)
          df[Boostername]

Out[52]:
```

	Date	Time (UTC)	Booster_Version	Launch_Site	Payload	PAYLOAD_MASS_KG_	Orbit	Customer	Mission_Outcome	Landing_Outcome
32	01-05-2017	11:15:00	F9 FT B1032.1	KSC LC-39A	NROL-76	5300	LEO	NRO	Success	Success (ground pad)
40	07-09-2017	14:00:00	F9 B4 B1040.1	KSC LC-39A	Boeing X-37B OTV-5	4990	LEO	U.S. Air Force	Success	Success (ground pad)
46	08-01-2018	01:00:00	F9 B4 B1043.1	CCAFS SLC-40	Zuma	5000	LEO	Northrop Grumman	Success (payload status unclear)	Success (ground pad)

```
In [53]: df[Boostername]["Booster_Version"]

Out[53]: 32      F9 FT B1032.1
          40      F9 B4 B1040.1
          46      F9 B4 B1043.1
          Name: Booster_Version, dtype: object
```

- SQLite Query Word: **AND(&)**
- The clause == filters for boosters which successfully landed on a [drone ship]
- The & condition determines successful landings with certain payload mass in the desired range
- Three F9 rockets met this condition

# Total Number of Successful and Failure Mission Outcomes

```
List the total number of successful and failure mission outcomes

In [57]: df["Mission_Outcome"].count()
#Total Mission Outcomes = 101

Out[57]: 101

In [66]: #Successful Mission Outcomes
Suc = (df["Mission_Outcome"] == "Success") | (df["Mission_Outcome"] == 'Success (payload status unclear)') |
df[Suc]
df[Suc].count()

Out[66]: Date          100
Time (UTC)         100
Booster_Version    100
Launch_Site        100
Payload            100
PAYLOAD_MASS_KG_   100
Orbit              100
Customer           100
Mission_Outcome     100
Landing_Outcome     100
dtype: int64
```

```
In [65]: df["Mission_Outcome"] == "Failure (in flight)"
Fai = (df["Mission_Outcome"] == "Failure (in flight)")
df[Fai]
df[Fai].count()

Out[65]: Date          1
Time (UTC)         1
Booster_Version    1
Launch_Site        1
Payload            1
PAYLOAD_MASS_KG_   1
Orbit              1
Customer           1
Mission_Outcome     1
Landing_Outcome     1
dtype: int64

In [64]: df["Mission_Outcome"].unique()

Out[64]: array(['Success', 'Failure (in flight)',
               'Success (payload status unclear)', 'Success '], dtype=object)
```

- SQLite Query Words: COUNT, UNIQUE
- 100 successful outcomes, 1 failed outcome
- This query identified specific parameters for successful or failed mission outcomes, including one qualified category (success, but payload status was unclear).

# Boosters Carried Maximum Payload

List the names of the booster\_versions which have carried the maximum payload mass. Use a subquery

```
In [75]: MaxVal = df["PAYLOAD_MASS_KG"].max()
df["PAYLOAD_MASS_KG"] == MaxVal
bol = (df["PAYLOAD_MASS_KG"] == MaxVal)
df[bol]["Booster_Version"]
```

```
Out[75]: 74      F9 B5 B1048.4
77      F9 B5 B1049.4
79      F9 B5 B1051.3
80      F9 B5 B1056.4
82      F9 B5 B1048.5
83      F9 B5 B1051.4
85      F9 B5 B1049.5
92      F9 B5 B1060.2
93      F9 B5 B1058.3
94      F9 B5 B1051.6
95      F9 B5 B1060.3
99      F9 B5 B1049.7
Name: Booster_Version, dtype: object
```

- SQLite Query Words: == clause, and the MAX() function
- The booster which carries the maximum payload is the F9 B5 version



# 2017 Launch Records

In [177...

```
df2 = df
df2['Year'] = df['Date'].apply(lambda x: int(x.split('-')[2]))
df2['Month'] = df['Date'].apply(lambda x: int(x.split('-')[1]))

b = (df2['Year']==2017) & (df2['Landing_Outcome']=='Success (ground pad)')
df2[b].loc[:, ['Month', 'Booster_Version', 'Launch_Site', 'Landing_Outcome']]
```

Out[177...

	Month	Booster_Version	Launch_Site	Landing_Outcome
29	2	F9 FT B1031.1	KSC LC-39A	Success (ground pad)
32	5	F9 FT B1032.1	KSC LC-39A	Success (ground pad)
34	6	F9 FT B1035.1	KSC LC-39A	Success (ground pad)
38	8	F9 B4 B1039.1	KSC LC-39A	Success (ground pad)
40	9	F9 B4 B1040.1	KSC LC-39A	Success (ground pad)
44	12	F9 FT B1035.2	CCAFS SLC-40	Success (ground pad)

- Due to utilizing SQLite, which does not display months, I had to separate and reclassify certain columns using Lambda as a window function; months are displayed numerically

# Rank Landing Outcomes Between 2010-06-04 and 2017-03-20

```
Rank the count of successful landing_outcomes between the date 04-06-2010 and 20-03-2017 in descending order.
```

```
In [201...] df['Landing_Outcome'].unique()
```

```
Out[201...] array(['Failure (parachute)', 'No attempt', 'Uncontrolled (ocean)',  
      'Controlled (ocean)', 'Failure (drone ship)',  
      'Precluded (drone ship)', 'Success (ground pad)',  
      'Success (drone ship)', 'Success', 'Failure', 'No attempt '],  
      dtype=object)
```

```
In [213...] %sql select "landing_outcome", count(*) from (select * from SPACEXTBL where (date between '04-06-2010' and  
* sqlite:///my_data1.db  
Done.
```

```
Out[213...]  Landing_Outcome  count(*)  
          Success          20  
    Success (drone ship)          8  
    Success (ground pad)          6
```

- SQ Lite Query Words: COUNT, WHERE, BETWEEN – used to select the correct of landing outcomes in the desired range.
- Additionally, the outcomes were SELECTED from **SpaceXTBL** (failure, success), to show descending order



# Section 3

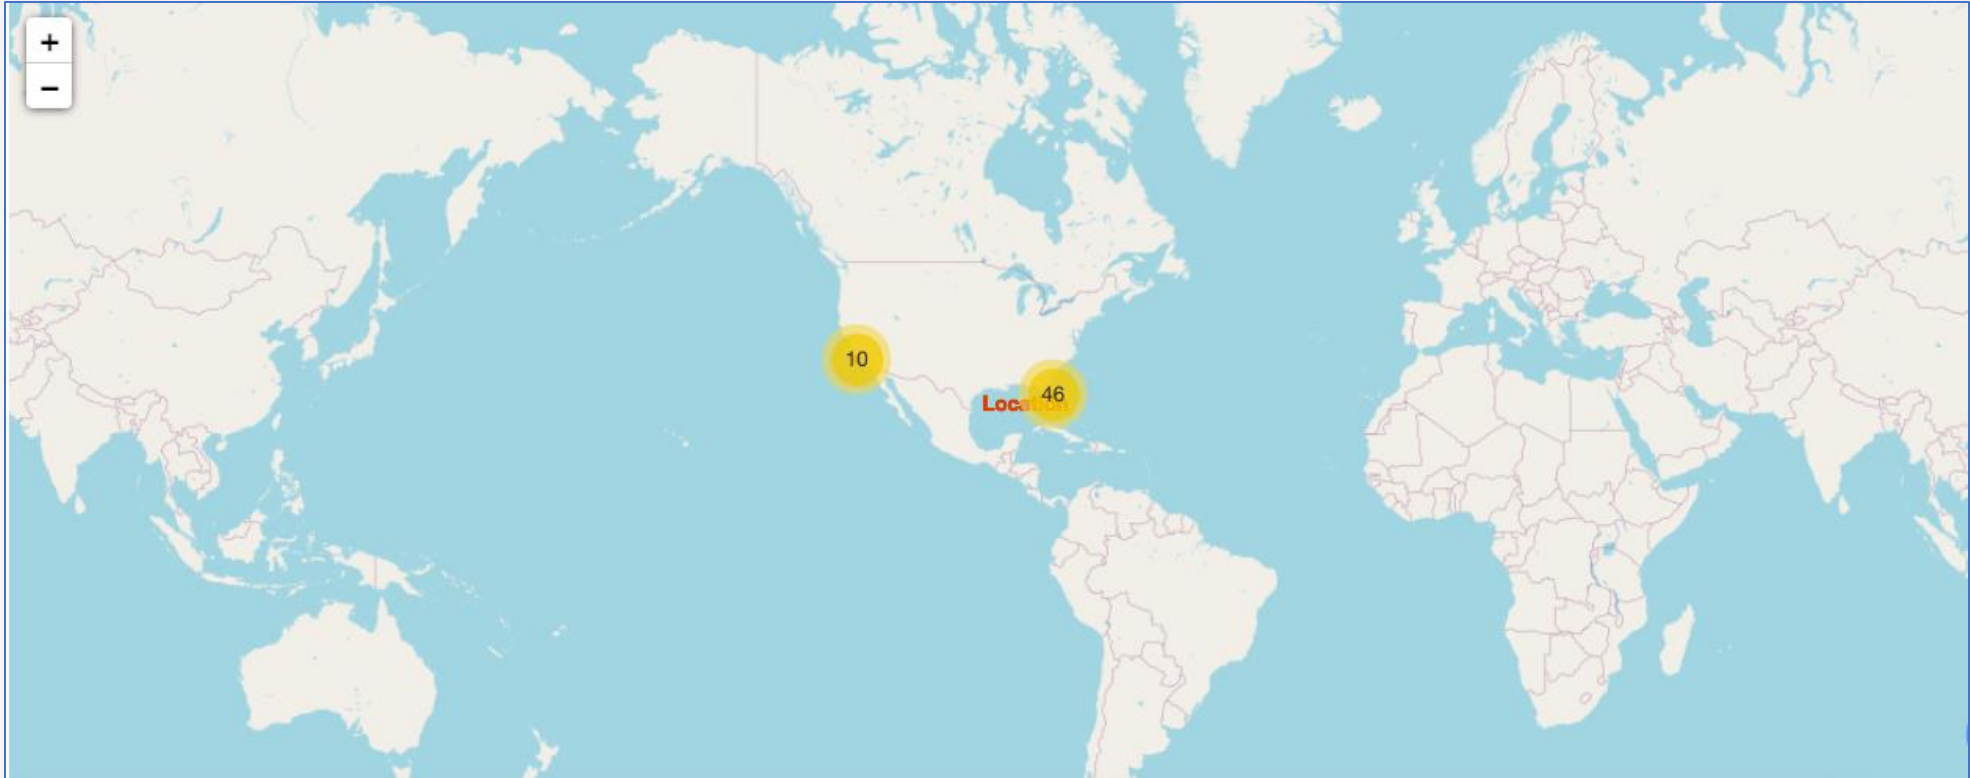
---

## Launch Sites Proximities Analysis





# All Launch Sites – Global Map Markers

---



Launch Sites for SpaceX Falcon 9 rockets are in California, on the **west coast** of the United States (Vandenberg AFB), and in Florida, on the **east coast** of the United States (Cape Canaveral).

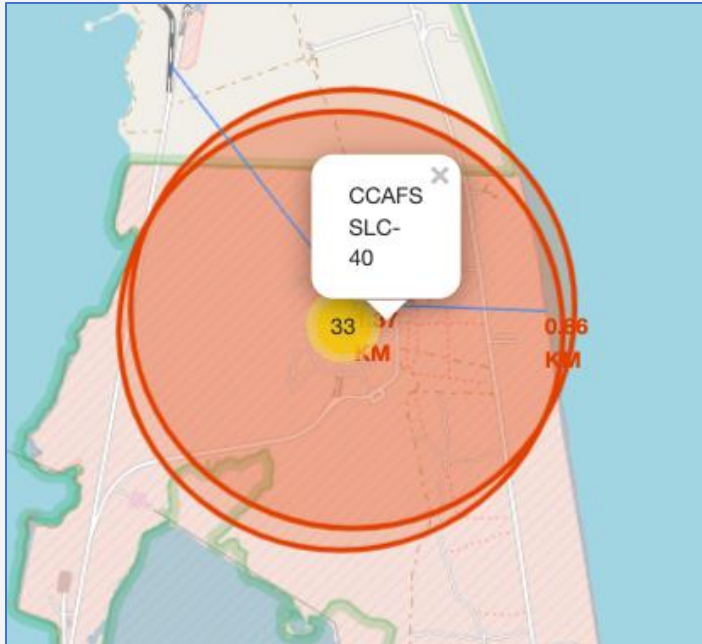




Markers Showing Launch Sites with Color Labels

A **GREEN** marker indicates a **successful** launch; whereas a **RED** marker indicates a **failed** launch

# Launch Site Distances to Landmarks



Distance from the Florida coastline



Distance to the closest Florida highway



Distance to the closet metropolitan city center

The Interactive Folium maps showed:

- Launch sites are NOT close to highways or metro city centers
- Railway proximities depend on private use (Space Force); typically commercial railways are avoided
- SpaceX certainly launch rockets close to coastlines






# Section 4

---

Building a Dashboard  
with Plotly Dash




# Dashboard: Pie Chart of All Successful Launch Sites



KSC LC-39A was the most successful launch site, with a rate of over 40%.




## Dashboard: Pie Chart for the Launch Site with the Highest Launch Success Ratio



KSC LC-39A achieved a 76.9% success rate for rocket launches. Thus, the failure rate stands at 23.1% at this launch site.

# Dashboard: Payload vs. Launch Scatterplot Chart for All Sites, with Slider Ranges



Lighter Payload:  
Slider Scale 2500 – 7000kg



Heavy Payload:  
Slider Scale 0 – 10,000kg

Success rates for lighter payloads are **higher** than heavyweight payloads.



## Section 5

---

### Predictive Analysis (Classification)



# Classification Accuracy

KNN, Decision Tree, and Logistic Regression all performed strongly – but numerically, the model with the highest numerical accuracy is the **Decision Tree**.


```
models = {'KNeighbors':knn_cv.best_score_,
          'DecisionTree':tree_cv.best_score_,
          'LogisticRegression':logreg_cv.best_score_,
          'SupportVector': svm_cv.best_score_}

bestalgorithm = max(models, key=models.get)
print('Best model is', bestalgorithm, 'with a score of', models[bestalgorithm])
if bestalgorithm == 'DecisionTree':
    print('Best params is:', tree_cv.best_params_)
if bestalgorithm == 'KNeighbors':
    print('Best params is:', knn_cv.best_params_)
if bestalgorithm == 'LogisticRegression':
    print('Best params is:', logreg_cv.best_params_)
if bestalgorithm == 'SupportVector':
    print('Best params is:', svm_cv.best_params_)

Best model is DecisionTree with a score of 0.8732142857142856
Best params is : {'criterion': 'gini', 'max_depth': 6, 'max_features': 'auto', 'min_samples_leaf': 2, 'min_samples_split': 5, 'splitter': 'random'}
```

```
In [37]: import numpy as np
import matplotlib.pyplot as plt

plt.barh(methods, accuracy)
plt.xlabel("Accuracy Percentage")
plt.ylabel("Method of Analysis")
plt.show()
```




After selecting the best hyperparameters for the Decision Tree classifier using the validation data, the model achieved 83.33% accuracy on the test data.

# Confusion Matrix

The Decision Tree machine learning model is able to distinguish between the different classes.

The major issue is the number of **false positives**, where failed launches are marked as successful. This information has a direct impact on deciding whether SpaceX launches can be accurately predicted.



# Case Study: SpaceX Rocket Launch

## Data Driven Insights & Conclusions

GEO, HEO, SSO, ES-L1 orbits have the best success rate



Low weight payloads perform better than heavier ones

SpaceX launches are proportionally successful; better performance is commensurate with time since the first launch

KSC LC-39A was the most successful launch site

The Decision Tree classifier is the best Machine Learning algorithm

# Appendix

Photograph credits for section headings are all sourced from commercial-free, royalty-free, open-source wallpaper websites:

- [wallpapercave.com](https://wallpapercave.com)
- [unsplash.com](https://unsplash.com)
- [pexels.com](https://pexels.com)



"Data will talk if you're willing to listen."

- Jim Bergeson